

MemoKi: Results and Validation Report (Gaps 1 to 4)

October 2026. What was built, what each experiment found, and how well each claim is validated. Hindi summary at the end.

Status labels used throughout

- **Supported** = measured, with evidence strong enough for a development-stage claim
- **Partly supported** = measured, but with a caveat that limits the claim
- **Not supported** = measured, and the data does not back the claim
- **Not yet tested** = no experiment has been run

Scope warning: all results are on the **validation split, 100 queries, one seed (11)**. They are development results. Final claims need the test split, several seeds and confidence intervals (section 8).

1. Summary

MemoKi stores whole conversations in a balanced tree whose nodes carry cheap summaries (names and date spans). A query uses cheap cue checks to skip branches before any vector comparison, with **zero LLM calls**.

Question	Answer so far
Does it beat brute-force search?	Yes: equal or better recall at 3% to 16% of the comparisons
Do cues really help?	Yes, significantly (p = 0.012 against the same tree without cues)
Does the hierarchy beat the best flat method?	No . A metadata filter plus HNSW is equal or better above about 230 comparisons. The tree wins only at very low budgets
Is it faster?	Not shown . Cue extraction costs about 5 ms per query
Can misfiled conversations be repaired?	Yes, by an offline audit (92% of lost queries recovered). Query-time fallback fails
Outdated facts (Gap 3)?	Not yet tested
Comparison with Mem0 and HiMem (Gap 4)?	Not yet done

2. Experimental setup

Item	Value
Data	LongMemEval-S: 19,195 conversation sessions, 470 answerable questions. LoCoMo downloaded for the Mem0/HiMem comparison (not used yet)
Split	Questions split 141 validation / 329 test, seeded and saved. Design decisions use validation only

validation only

Query set	100 validation queries (up to 2 evidence sessions each), the same at every store size
Store sizes	1,000, 5,000, 10,000 (also 500 in early runs). Smaller stores are subsets of larger ones
Embeddings	all-MiniLM-L6-v2, 256 tokens, computed once, shared by every method
Metric	Recall@5 (a query is a hit if any relevant conversation is in the top 5)
Cost	Embedding comparisons per query (cue checks counted separately). Latency is secondary
Hardware	CPU laptop, 8 cores, 7.8 GB RAM, no usable GPU
Tests	64 automated tests pass

3. Methods compared

Method	What it is
Brute force, FAISS flat	Compare the query with every conversation
HNSW (ef 5 to 128)	Standard approximate nearest-neighbour index
Flat + cue filter	Flat store with the same name and date filter via an inverted index
Filter + HNSW	The strongest flat opponent: cue filter when cues exist, HNSW otherwise
Tree, no cues	Our tree with centroid routing only (32 children per node, height 3 at 10,000)
Cue tree (proposed)	Same tree with name and date summaries on every node, cue gate before vector comparison
Cue tree + audit	Cue tree after one offline re-filing pass

4. Claims and their validation

#	Claim	Evidence	Status	To confirm
C1	Cue tree matches brute-force accuracy at a small share of the cost	Recall 0.68 / 0.47 / 0.42 vs 0.63 / 0.46 / 0.36, using 16% / 6% / 3% of comparisons	Supported (dev)	Test split, seeds, CIs
C2	Cues significantly improve the tree	10 wins, 1 loss vs the same tree without cues, $p = 0.012$	Supported	Test split
C3	Cue tree beats brute force on queries with cues	0.634 vs 0.488 (41 queries); 8 wins vs 2 losses, $p = 0.109$	Partly supported	More queries

C4	Cue tree beats flat search including HNSW	Hybrid filter + HNSW: 0.440 at 297 comps vs 0.420 at 308 (10,000 sessions); 2 vs 4 queries, $p = 0.69$	Not supported	Cannot be claimed
C5	Cue tree is better at very low budgets	0.350 at 108 comparisons; HNSW needs about 320 for 0.350	Partly supported	Test split; this is the tree's niche
C6	Cost grows slowly with store size	Cue tree 162, 303, 308 comparisons at 1k, 5k, 10k. HNSW also flat (255, 316, 320)	Supported , but not unique to the tree	50k to 100k stress test
C7	Faster than flat search	Search 0.27 ms vs 0.98 ms (brute force) at 10k, but query prep about 5 ms, HNSW 0.05 to 0.15 ms	Not supported end to end	Cheaper cue extractor
C8	Zero LLM calls	By construction; counter is 0 everywhere	Supported	None
C9	Misfiling costs recall	-0.2, -2.0, -3.8 points at 5%, 10%, 20%	Supported (small counts)	Test split, 10k
C10	Offline audit repairs misfiling	3/4, 11/12, 23/25 lost queries recovered at +2% query cost	Partly supported	Test split, other sizes, margin check
C11	Query-time flat fallback repairs misfiling	At the synopsis criterion: no setting passes	Not supported	Reported as a negative result
C12	Audit also improves a healthy tree	+2 to +4 pts at 5k, but -1 pt at 1k and -2 pts at 10k	Not supported	Treated as noise
C13	Entity cue is reliable	Recall 95% (37/39), selectivity mean 1.2%	Partly supported	Rerun on validation; label sheet
C14	Time cue is reliable but rare	Applies to 8% of queries; recall 11/11	Partly supported	More queries
C15	Recency rule handles outdated facts	No experiment	Not yet tested	Gap 3
C16	Better than Mem0 and HiMem on LLM calls, latency, tokens	No experiment	Not yet tested	Gap 4

5. Results by gap

Gap 1: cost against scale

Validation, beam 10 for the tree. Recall@5 and mean embedding comparisons per query:

Method	1,000 sessions	5,000 sessions	10,000 sessions
Brute force	0.630 / 1,000	0.460 / 5,000	0.360 / 10,000
HNSW (ef 16)	0.640 / 255	0.450 / 316	0.350 / 320
Filter + HNSW (ef 16)	0.710 / 162	0.520 / 243	0.440 / 297
Tree, no cues	0.600 / 254	0.390 / 463	0.330 / 457
Cue tree	0.680 / 162	0.470 / 303	0.420 / 308

Other findings:

- **Low-budget regime (10,000 sessions):** cue tree at beam 3 gives 0.350 with 108 comparisons. HNSW at ef 5 gives 0.310 with 201, and at ef 8 gives 0.330 with 237.
- **Queries with a cue (41) vs without (59), 10,000 sessions, beam 10:** brute force 0.488 / 0.271; tree without cues 0.415 / 0.271; cue tree 0.634 / 0.271. Queries without cues are unchanged by design.
- **Tree structure:** frozen at 32 children and minimum fill 0.35 by a rule written before the sweep. At 10,000 sessions the height fell from 7 to 3 and the share of brute-force hits kept rose from 61% to 86%.
- **Cue quality:** entities appear in 36% of queries (recall 95%, one name check removes about 99% of distractors). Dates apply to 8% of queries (recall 11 of 11, 13.7% of the pool left in the window).
- **Ablation, beam 10 (no cues / names / names + time):** 1,000 sessions 0.600 / 0.670 / 0.680; 5,000 sessions 0.390 / 0.460 / 0.470; 10,000 sessions 0.330 / 0.410 / 0.420. Time adds about 1 point.
- **Time:** search 0.27 ms against 0.98 ms (brute force) at 10,000 sessions, but query preparation is about 5 ms, so end to end the cue tree is about 4.9 ms against 0.06 to 0.15 ms for HNSW.

Verdict: partly supported. Strong against brute force; a tie or slight loss against filter + HNSW; no latency claim.

Gap 2: misfiled conversations

Damage (5,000 sessions, beam 10, 5 repeats):

Misfiled	Recall@5 (clean 0.470)	Lost queries	Cue / no-cue
5%	0.468	4	0 / 4
10%	0.450	12	4 / 8
20%	0.432	25	6 / 19

Cue summaries protect partly: cue queries are 41% of queries but 24 to 33% of the losses.

Query-time recovery, 10% misfiling (best settings per tier):

Tier	Recovery	Extra cost	Note
------	----------	------------	------

Entity lookup, always	4/12 (33%)	+5.1%	NO harm
Wider beam x8, always	10/12 (83%)	+397%	Equivalent to a bigger beam
Full flat, always	12/12 (100%)	+1,642%	Breaks 26 correct queries
Full flat, threshold 0.30	2/12 (17%)	+177%	Breaks 15 correct queries

No tier meets the synopsis criterion (over 80% recovery with under 25% extra). The "weak similarity" trigger does not separate lost queries from normal ones.

Offline audit (re-file a conversation if another group is closer):

Misfiled	Lost	Recovered	Query-time comparisons
5%	4	3 (75%)	303 to 310
10%	12	11 (92%)	303 to 311
20%	25	23 (92%)	303 to 312

It finds 98 to 100% of the deliberately misfiled leaves and meets the criterion. But my misfiling is easy to detect (a far-away branch), and the audit on a healthy tree gave inconsistent results: cue tree recall 0.680 to 0.670 at 1,000, 0.470 to 0.490 at 5,000, 0.420 to 0.400 at 10,000 (none significant).

Verdict: largely answered on development data. Still to do: the cue-accuracy label sheet (not filled in), the audit margin check at other sizes, and the test split.

Gap 3: outdated information

Not yet tested. Ready: the date parser, timestamps on every node, and 72 knowledge-update queries. To do: build 50 to 100 update pairs, add a recency bonus that applies when similarities are close, tune its two parameters on validation, and measure the stale-retrieval rate and freshness accuracy against the Recall@5 cost (target: at most 1 to 2 points).

Gap 4: fair comparison

Done	Not done
One config, same data, queries, embedding model, seeds and split for every method	Mem0 and HiMem on LoCoMo (LLM calls, latency, retrieved tokens)
Shared cost counters; cost counted honestly (cue checks, prep time)	Type and emotion cues (not built)
Rules written before results; paired exact tests	Repeats over 3 to 5 seeds with mean and standard deviation
Strong opponents (HNSW, filter + HNSW)	Full ablation table and the test split
Partial ablation: no cues / names / names + time / fallback	

The HiMem paper states its code is at github.com/jojopdq/HiMem. I have not checked how hard it is to run or what LLM it needs. Your laptop (CPU only, 7.8 GB RAM) will likely limit this to a small slice of LoCoMo, with extrapolated costs stated as such.

Verdict: foundation solid, main comparisons ahead.

6. Validation hygiene

What protects the results

- Fixed, saved validation and test split; no design decisions on the test split (with exceptions below).
- Rules written before results: tree structure, fallback threshold, audit margin.
- Paired exact tests, not just averages. Counts, not milliseconds, as the primary cost measure.
- Strong opponents included. Five bugs found and fixed (cache loss after a crash, a date-phrase crash, a stale cue cache, negative search time, model-load inflation of prep time).

What weakens the results (state these in the final report)

1. **100 queries, one seed.** One query is 1 point. Many differences are 1 to 4 queries.
 2. **Small denominators:** Gap 2 lost counts are 4, 12 and 25, and repeats reuse the same queries, so they are correlated.
 3. **Early checks used test queries:** the first baseline table, the model comparison and the first entity check. They informed only choices that were then kept as defaults (MiniLM, entity cue), but they must be disclosed. The entity check still needs a validation rerun.
 4. **One rule was amended after seeing data:** the audit margin must satisfy its clean-tree criterion at every size, not only at 5,000. The amendment only makes the choice stricter.
 5. **My misfiling simulation is easier than natural misfiling.**
 6. **Timing noise:** sub-millisecond timings vary between runs on a laptop.
 7. **Cue extraction was slower than search,** so no latency claim is made.
-

7. What we can say now, and what we cannot

Defensible now (development evidence)

- A cue-guided hierarchical memory reaches brute-force accuracy at about 3% to 16% of the comparisons, with no LLM calls.
- Cues give a statistically significant gain over the same hierarchy without them ($p = 0.012$).
- On queries with cues, recall is far above brute force (0.634 against 0.488), though this effect is shared with a flat metadata filter.
- Misfiling costs 2 to 4 recall points at 10% to 20% misfiled, and an offline audit repairs most of it at almost no query-time cost.

Not defensible

- "Faster than flat search" or "beats HNSW".
 - "The hierarchy is why accuracy improves" (a flat filter gets the same gain).
 - Any statement about Gap 3, or about Mem0 and HiMem.
 - Anything stated as final before the test-split run.
-

8. What remains, and how each item validates a claim

Step	Validates	Output
Refile check at 1,000 and 10,000 sessions; freeze audit margin	C10, C12	results\refile_1000_10000.txt
Entity check rerun on validation split	C13	rerun eval.cue_selectivity
Hand-label 100 conversations (eval.label_sample, eval.score_labels)	C13, Gap 2 cue accuracy	precision, recall of names
Recency rule and update-pair experiment	C15 (Gap 3)	stale rate, freshness accuracy
Ollama timing test, then Mem0/HiMem spike	C16 (Gap 4)	LLM calls, tokens, time per session
Type and emotion cues, tested like the others	Gap 4 ablation	recall and selectivity
Final test-split run: 3 to 5 seeds, all five sizes, more queries, confidence intervals	C1 to C14	final tables
50k to 100k stress test	C6, C7	scaling beyond 10k
Server, extension, dashboard	Delivery	working demo

9. Reproducing each result

Claim	Command	Saved in results\
Tests	pytest	64 passed
Tree structure	python -m eval.tree_sweep	console output
Cue tree vs no cues, fallback	python -m eval.cue_tree_check (and -taus 0.30 0.40)	cue_tree_validation.txt, cue_tree_fallback_validation.txt
Every method, side by side	python -m eval.compare_check --sizes 1000 5000 10000 --efs 5 8 16 32 64 128 --beams 1 2 3 5 10 20 --audit-margin 0.0 --csv results\pareto_audit.csv	pareto_audit_table.txt, pareto_audit.csv, pareto_*.png
Misfiling and query-time recovery	python -m eval.misfile_check --sizes 5000	misfile_5000.txt
Offline	python -m eval.refile_check --sizes	refile_5000.txt

audit 5000
repair

Date cue python -m eval.time_cue_check console output

10/10/2026